Governing Power BI & Fabric

When AI Writes the Reports

Pragmatic governance strategies for AI-assisted analytics environments

Steve Fox
Principal Consultant : Governance
AnalyticsinAction.com

This space moves fast. Content reflects thinking as of January 2026.

The Agent Explosion

AI-assisted development has reached a tipping point

- Explosive growth in agentic AI tools for Power BI and Fabric
- Created by: Microsoft, vendors, community projects, individuals
- Capabilities: Create semantic models, write DAX, automate CI/CD
- Shadow use is high - users aren't waiting for IT approval
- Diverse architectures

VENDOR

Copilot, Fabric Data Agents etc.

COMMUNITY

Tabular Editor MCP , Fabric MCP

PERSONAL TOOLS

Personal Subscription Claude, ChatGPT + MCP

If you're not already seeing agents on your tenant, its probably because tracking AI usage is nearly impossible.

Not All Risks Are Equal

Risk depends on where your data ends up

LOWER RISK	MEDIUM	HIGHER RISK
Local LLMs (reputable vendors) Amazon Bedrock Fabric Data Agents	Azure OpenAI Enterprise APIs (with opt-outs)	ChatGPT (consumer) Claude (consumer)
Data stays local No training exposure	Contractual protection	May feed training sets

Other Risk Dimensions

- Probabilistic outputs | Naive by nature | External connections (MCPs etc.) | Inherited privileges | Leaky guardrails | Autonomy | human oversight and falling asleep at the wheel.

NY Times vs OpenAI lawsuit: Copyright content in Open AI's training sets

Specific Risks

OWASP LLM Top 10 (2025) framework

OWASP LLM 2025

LLM01	Prompt Injection	LLM06	**Excessive Agency**
LLM02	**Sensitive Info Disclosure**	LLM07	System Prompt Leakage
LLM03	Supply Chain	LLM08	Vector/Embedding Weakness
LLM04	Data and Model Poisoning	LLM09	**Misinformation**
LLM05	Improper Output Handling	LLM10	**Unbounded Consumption**

TODAY'S FOCUS

- Deletion (Excessive Agency)
- Data Leakage (Info Disclosure)
- Slop Generation (Misinformation)
- Unbounded Consumption

Arguably these 4 risks are the most relevant to Power BI and Fabric governance today.

Four Risks to Keep Top of Mind

Pragmatic Governance: Encourage safer use, not block entirely

Excessive agency (Data deletion)

Production data or content wiped by agents with write access

Catastrophic, visible

Misinformation (Slop Generation)

AI creates content faster than you can review or govern it

Gradual, trust-eroding

Sensitive Information Disclosure

Sensitive information sent to external LLM services

Potentially invisible

Unbounded Consumption

Agent burns through Fabric capacity with large context queries

Costly, disruptive

Key question: Is the change reversible? Reversibility determines severity.

The Choice

Two paths forward

PATH A: BLOCK

- Ban all AI tools
- Lock down tenant
- Heavy restrictions

RESULT:

- Shadow IT increases
- Less safe usage
- Talent leaves
- Fall behind

PATH B: ENABLE

- Provide approved alternatives
- Educate on risks
- Implement guardrails

RESULT:

- Controlled adoption
- Visibility improves
- Productivity gains
- Governance matures

Risk: If you don't provide options, users will find their own - and those options will be less safe.

Controls

Educate, Isolate, limit privileges, ensure recoverability

EDUCATE

- Ensure users are aware of risks
- Encourage and monitor best practices
- Develop internal AI Analyst user groups
- Regular training to keep up with AI changes / Tools
- Skill development to mitigate risks (CI/CD etc.)

ISOLATION

- Keep agents away from production
- Separate dev, test, and production environments
- Consider capacity isolation (dedicated Fabric capacity)
- Sandbox options: containers, virtual desktops

PRIVILEGE

- Agents often inherit user privileges - this is dangerous
- User with tenant admin = Agent with tenant admin
- Consider feasibility of agent-specific identities
- Assign minimal permissions: Contributor, not Admin

RECOVERY

- If an agent breaks something, can you recover?
- Power BI now supports 5 version reverts
- Implement CI/CD pipelines
- Source control everything (Git integration / SharePoint)

Automated Governance

Leverage AI to help govern what AI creates

Jevons Paradox: When something becomes cheaper/easier to produce, we produce MORE of it, not less. AI-generated content will explode in volume - manual review cannot scale.

AUTOMATED TESTING

- Use AI to generate tests
- Integrate testing into CI/CD & pipelines
- Run BPA programmatically to avoid Power BI slop
- Performance test before deploy

VALIDATION GATES

- Separate potentially destructive AI from production
- Human review checkpoints
- Automated quality thresholds
- Performance and compliance gates

HUMAN STANDARDS

- Same testing as human code
- Same review process
- Same documentation standards
- "AI-generated" is not an excuse

Governance Pipeline:

AI Creates → Auto Tests → Validation → Review → Production

Testing and oversight becomes MORE important with AI, not less.

What To Do Now

Four actions you can take immediately

Three Control Areas

Human Controls
Educate, policy, accountability

Automated Governance
Test, validate, gate

Technical Controls
Isolate, manage, privileges, backup,
Drive use of safer tools

START MONDAY

1. Draft & socialize an AI usage policy (include “why”)
2. Audit role-based access to workspaces
3. Explore one isolation pattern (separate dev from prod)
4. Short list AI tools to provision for your team

The gate is open – its will be easier to act before adoption accelerates

Resources and Deeper Dives

Scan to access all references

What You'll Find

- Slides (PDF download)
- Referenced articles
- How to reach me : linkedin.com/in/stevefoxnz/

Thank You!

Questions?

Global Power BI and Fabric Summit 2026
23-27 February 2026